

03 Exercise Notebook 3

March 29, 2023

1 Exercise 3

In this exercise, you will analyse a dataset obtained from the London transport system (TfL). The data is in a file called `tfl_readership.csv` (comma-separated-values format). As in Exercise 2, we will load and view the data using `pandas`.

```
[1]: # If you are running this on Google Colab, uncomment and run the following ↵  
      ↪ lines; otherwise ignore this cell  
      # from google.colab import drive  
      # drive.mount('/content/drive')
```

```
[2]: import math  
      import numpy as np  
      import matplotlib.pyplot as plt  
      import pandas as pd
```

```
[3]: # Load data  
df_tfl = pd.read_csv('tfl_ridership.csv')  
# If running on Google Colab change path to '/content/drive/MyDrive/  
      ↪ IB-Data-Science/Exercises/tfl_ridership.csv'  
  
df_tfl.head(3)
```

```
[3]:
```

	Year	Period	Start	End	Days	Bus cash (000s)	\
0	2000/01	P 01	01 Apr '00	29 Apr '00	29d	884	
1	2000/01	P 02	30 Apr '00	27 May '00	28d	949	
2	2000/01	P 03	28 May '00	24 Jun '00	28d	945	

	Bus Oyster PAYG (000s)	Bus Contactless (000s)	\
0	0	0	
1	0	0	
2	0	0	

	Bus One Day Bus Pass (000s)	Bus Day Travelcard (000s)	...	\
0	210	231	...	
1	214	205	...	
2	209	221	...	

	Tube Contactless (000s)	Tube Day Travelcard (000s)	\
0	0	655	
1	0	605	
2	0	650	

	Tube Season Travelcard (000s)	Tube Other incl free (000s)	\
0	1066	200	
1	1168	217	
2	1154	212	

	Tube Total (000s)	TfL Rail (000s)	Overground (000s)	DLR (000s)	\
0	2509	0	0	96	
1	2598	0	0	93	
2	2623	0	0	98	

	Tram (000s)	Air Line (000s)
0	45.8	0.0
1	46.5	0.0
2	47.1	0.0

[3 rows x 26 columns]

Each row of our data frame represents the average daily ridership over a 28/29 day period for various types of transport and tickets (bus, tube etc.). We have used the `.head()` command to display the top 13 rows of the data frame (corresponding to one year). Focusing on the “Tube Total” column, notice the dip in ridership in row 9 (presumably due to Christmas/New Year’s), and also the slight dip during the summer (rows 4,5).

```
[4]: #df_tfl.sample(3) #random sample of 3 rows
df_tfl.tail(3) #last 3 rows
```

```
[4]:      Year Period      Start      End Days  Bus cash (000s) \
242  2018/19    P 09  11 Nov '18  08 Dec '18  28d           0
243  2018/19    P 10  09 Dec '18  05 Jan '19  28d           0
244  2018/19    P 11  06 Jan '19  02 Feb '19  28d           0
```

	Bus Oyster PAYG (000s)	Bus Contactless (000s)	\
242	1110	1089	
243	1001	949	
244	1036	1075	

	Bus One Day Bus Pass (000s)	Bus Day Travelcard (000s)	...	\
242	0	41	...	
243	0	38	...	
244	0	30	...	

	Tube Contactless (000s)	Tube Day Travelcard (000s)	\
242	1399	249	

243	1110	242
244	1310	204

	Tube Season Travelcard (000s)	Tube Other incl free (000s)	\
242	1017	334	
243	632	259	
244	924	305	

	Tube Total (000s)	TfL Rail (000s)	Overground (000s)	DLR (000s)	\
242	4221	996	557	355	
243	3279	750	414	270	
244	3809	929	517	333	

	Tram (000s)	Air Line (000s)
242	84.1	2.6
243	66.3	3.2
244	79.3	2.3

[3 rows x 26 columns]

The dataframe contains $N = 245$ counting periods (of 28/29 days each) from 1 April 2000 to 2 Feb 2019. We now define a numpy array consisting of the values in the 'Tube Total (000s)' column:

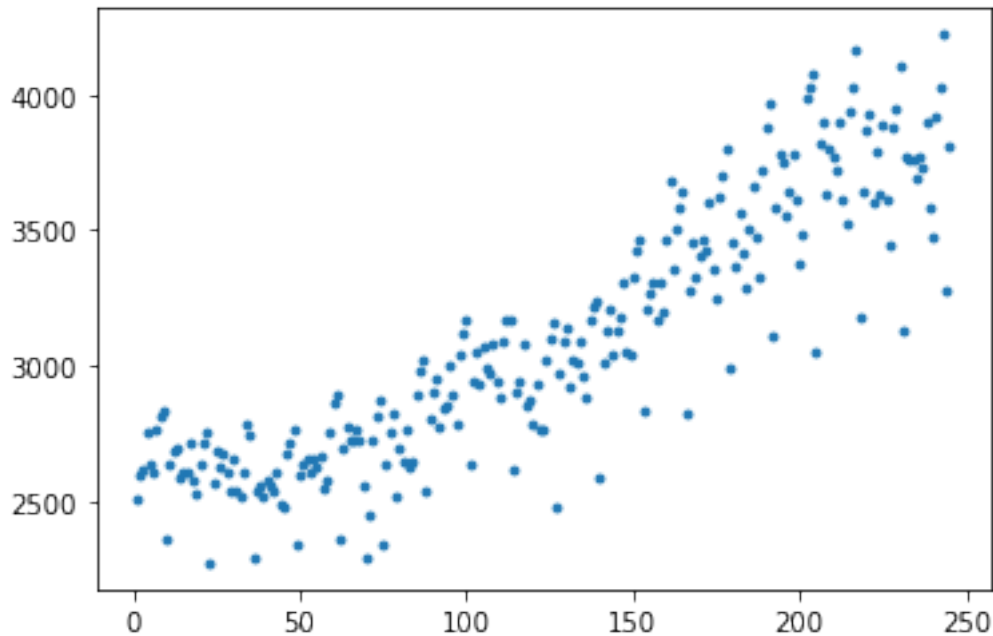
```
[5]: yvals = np.array(df_tfl['Tube Total (000s)'])
      N = np.size(yvals)
      xvals = np.linspace(1,N,N) #an array containing the values 1,2,...,N
```

We now have a time series consisting of points (x_i, y_i) , for $i = 1, \dots, N$, where y_i is the average daily tube rideship in counting period $x_i = i$.

1.1 2a) Plot the data in a scatterplot

```
[6]: #Your code for scatterplot here

plt.scatter(xvals, yvals, s=8)
plt.rcParams['figure.figsize'] = [15, 8]
plt.rcParams['axes.titlesize'] = 20
plt.rcParams['axes.labelsize'] = 20
plt.rcParams['xtick.labelsize'] = 14
plt.rcParams['ytick.labelsize'] = 14
plt.show()
```



1.2 2b) Fit a linear model $f(x) = \beta_0 + \beta_1 x$ to the data

- Print the values of the regression coefficients β_0, β_1 determined using least-squares.
- Plot the fitted model and the scatterplot on the same plot.
- Compute and print the **MSE** and the R^2 coefficient for the fitted model.

All numerical outputs should be displayed to three decimal places.

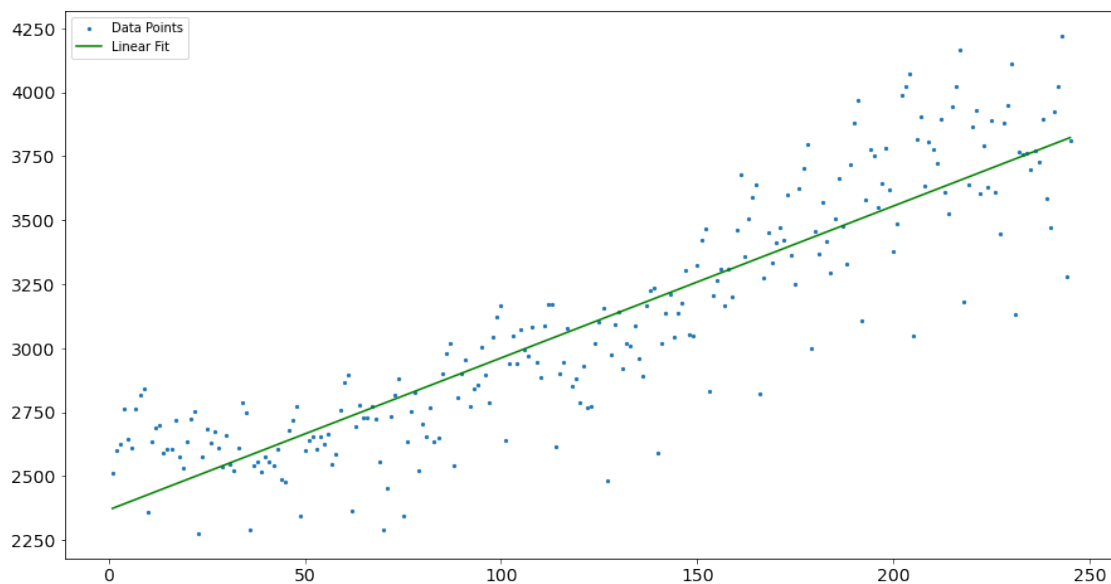
```
[7]: #Your code here
def polyreg(data_matrix: np.array, k: int)->np.array:
    """
    The function returns the the coefficient vector beta, the fit X*beta, and
    the vector of residuals y-X*beta
    """
    N, _ = data_matrix.shape
    assert _ == 2
    t, y = [data_matrix[:,i] for i in range(2)]
    X = np.ones((N,k+1))
    for i in range(1,k+1):
        X[:,i] = t**i
    P = np.linalg.inv((X.T).dot(X))
    P = P.dot(X.T)
    beta = P.dot(y)
    fit = np.dot(X, beta)
    residual = y - fit
    return (beta,fit,residual)
```

```

xy_data= np.column_stack([xvals, yvals])
beta, fit, residual = polyreg(xy_data, 1)
MSE = residual.dot(residual) /N
R2 = 1- (MSE/ yvals.var())
print(f'Least squares coefficients for k=1 linear regression: y =_
↳{round(beta[1],3)} x + {round(beta[0],3)}')
print(f'MSE1= {np.round(MSE,3)}, R2 = {np.round(R2,3)}')
plt.scatter(xvals, yvals, s=5, label='Data Points')
plt.plot(xvals, beta[0]+ beta[1]* xvals,'g',label='Linear Fit')
plt.legend()
plt.show()

```

Least squares coefficients for k=1 linear regression: $y = 5.939 x + 2367.382$
MSE1= 45323.636, R2 = 0.796



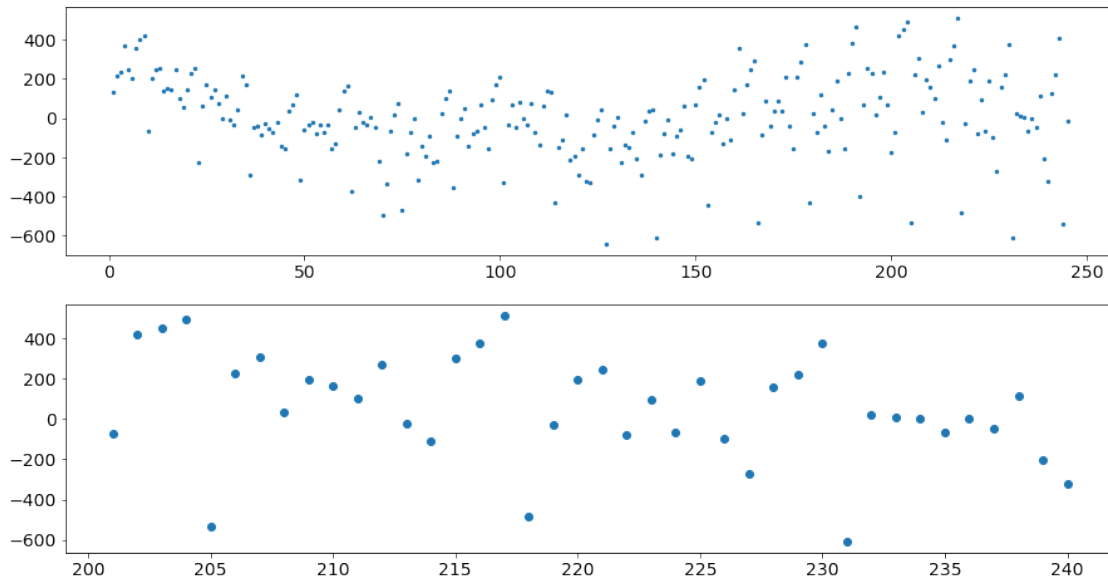
1.3 2c) Plotting the residuals

- Plot the residuals on a scatterplot
- Also plot the residuals over a short duration and comment on whether you can discern any periodic components.

```

[8]: # Your code here
plt.subplot(211)
plt.scatter(xvals, residual, s=6)
plt.subplot(212)
plt.scatter(xvals[200:240], residual[200:240])
plt.show()

```



Even though it is not a sinusoid, there is a repeating pattern whose period is about 13 time steps, which corresponds to about 364 days, so 1 year.

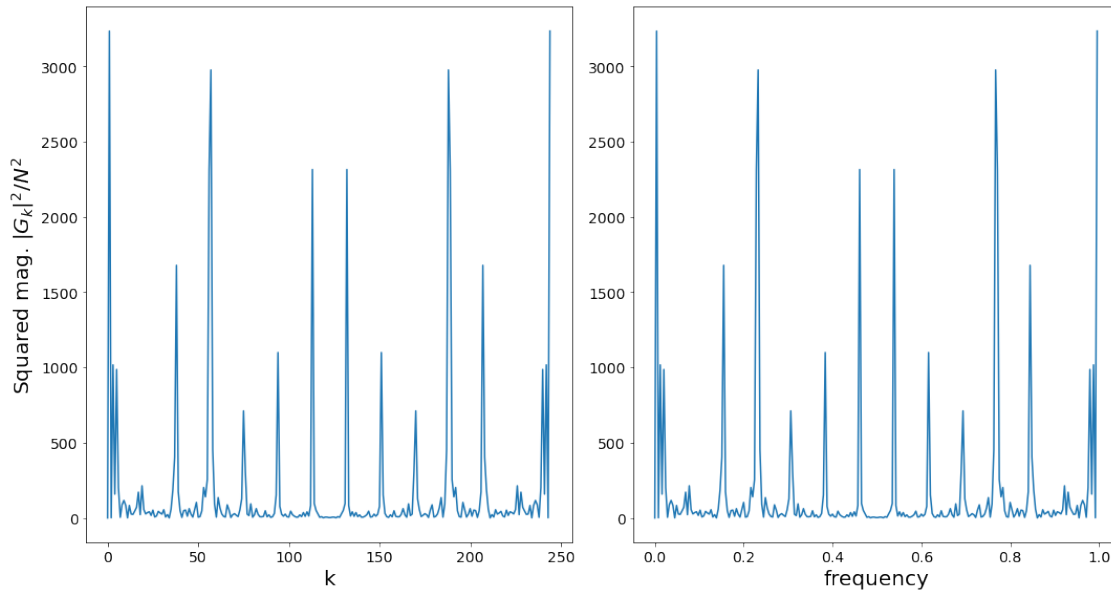
2d) Periodogram

- Compute and plot the periodogram of the residuals. (Recall that the periodogram is the squared-magnitude of the DFT coefficients.)
- Identify the indices/frequencies for which the periodogram value exceeds **50%** of the maximum.

```
[9]: # Your code to compute and plot the histogram
from numpy.fft import fft
df_tfl['days_parsed'] = df_tfl.Days.apply(lambda x: int(x[:-1]))
T = df_tfl.days_parsed.mean() / 365.5 ##Get the periodicity reported to 1 year
    ↳ instead of period count
residual_freq = fft(residual, N) / N
pdgram = np.abs(residual_freq)**2
indices = np.linspace(0, N-1, N)
freq = indices / N
angular_freq = 2 * np.pi * freq

plt.title('Periodogram of the linear fit residuals')
plt.subplot(121)
plt.plot(indices, pdgram)
plt.xlabel('k')
plt.ylabel('Squared mag.  $|G_k|^2/N^2$ ')
plt.subplot(122)
plt.plot(freq, pdgram)
plt.xlabel('frequency')
plt.tight_layout()
```

```
plt.show()
```



```
[10]: # Your code to identify the indices for which the periodogram value exceeds 50%
      ↪ of the maximum
high = pdgram.max()
top_ind = indices[(pdgram >= 0.5 * high)]
top_freq = freq[(pdgram >= 0.5 * high)]
assert top_ind.sum()%N==0
#Get rid of conjugate pairs
top_ind=top_ind[:top_ind.size//2]
top_freq=top_freq[:top_freq.size//2]

[11]: #Get temporal frequencies
time_top_freq = 1/T * top_freq
print(f'The most significant harmonic components correspond to the following
      ↪ frequencies in yr^-1 :\n{time_top_freq}')
```

The most significant harmonic components correspond to the following frequencies in yr^{-1} :

```
[0.05310956 2.01816332 2.97413543 3.02724499 6.00138041]
```

1.4 2e) To the residuals, fit a model of the form

$$\beta_{1s} \sin(\omega_1 x) + \beta_{1c} \cos(\omega_1 x) + \beta_{2s} \sin(\omega_2 x) + \beta_{2c} \cos(\omega_2 x) + \dots + \beta_{Ks} \sin(\omega_K x) + \beta_{Kc} \cos(\omega_K x).$$

The frequencies $\omega_1, \dots, \omega_K$ in the model are those corresponding to the indices identified in Part 2c. (Hint: Each of the sines and cosines will correspond to one column in your X-matrix.)

- Print the values of the regression coefficients obtained using least-squares.

- Compute and print the final **MSE** and R^2 coefficient. Comment on the improvement over the linear fit.

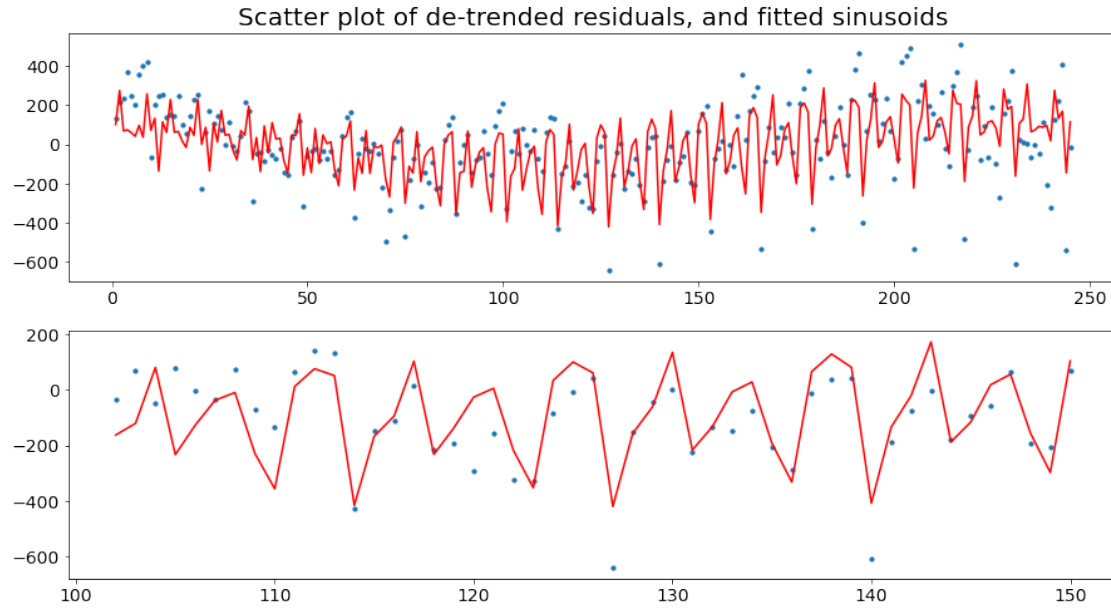
All numerical outputs should be displayed to three decimal places.

```
[12]: def harmonic_reg(t: np.array, y: np.array, frequencies: np.array)->np.array:
    '''
    The function returns the the coefficient vector beta, the fit X*beta, and
    the vector of residuals y-X*beta
    '''
    assert t.shape==y.shape
    N = t.size
    X = np.ones((N, 2* len(frequencies)))
    for i in range(len(frequencies)):
        X[:,2*i ] = np.sin(2 *np.pi *frequencies[i] *t)
        X[:,2*i+1] = np.cos(2 *np.pi *frequencies[i] *t)
    P = np.linalg.inv((X.T).dot(X))
    P = P.dot(X.T)
    beta = P.dot(y)
    fit = np.dot(X, beta)
    residual = y - fit
    return (beta,fit,residual)

beta_sc, fit_sc, residual_sc = harmonic_reg(xvals, residual, top_freq)
MSE_sc = residual_sc.dot(residual_sc) / residual_sc.size
R2_sc = 1 - MSE_sc / residual.var()
print(f'Least squares coefficients for linear regression:{np.round(beta_sc,3)}')
print(f'MSE_sc= {np.round(MSE_sc,3)}, R2 = {np.round(R2_sc,3)}')

plt.subplot(211)
plt.scatter(xvals, residual, s=10)
plt.plot(xvals, fit_sc, 'r')
plt.title('Scatter plot of de-trended residuals, and fitted sinusoids')
# Zoom in to a few values
plt.subplot(212)
plt.scatter(xvals[101:150], residual[101:150], s=10)
plt.plot(xvals[101:150], fit_sc[101:150], 'r')
plt.show()
```

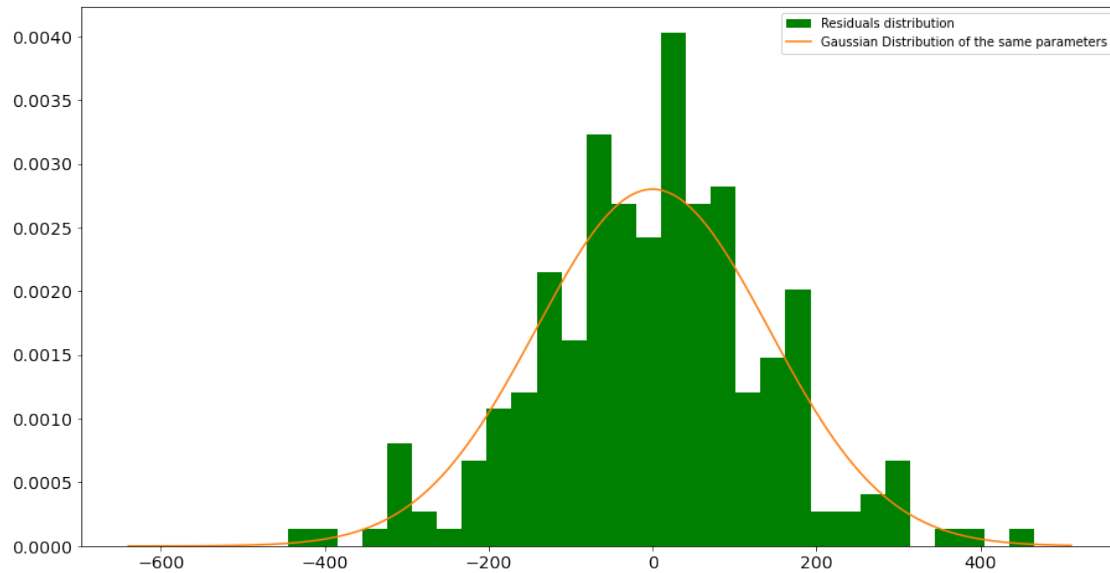
```
Least squares coefficients for linear regression: [-51.253 101.556 61.628
-54.006 -15.581 -94.797 81.659 72.381 32.472
90.589]
MSE_sc= 20297.501, R2 = 0.552
```

The MSE is halved by adding the seasonality correction, and a small R suggests that we can improve the precision by including more terms.

```
[13]: from scipy.stats import norm

n, bins, patches = plt.hist(residual_sc, bins=30, density=True,
    ↪facecolor='green',label='Residuals distribution')
x_pdf = np.linspace(residual.min(), residual.max(),1000)
y_pdf = norm.pdf(x_pdf,loc = residual_sc.mean(), scale = residual_sc.std() )
plt.plot(x_pdf,y_pdf,label='Gaussian Distribution of the same parameters')
plt.legend()
plt.show()
```



1.5 2f) The combined fit

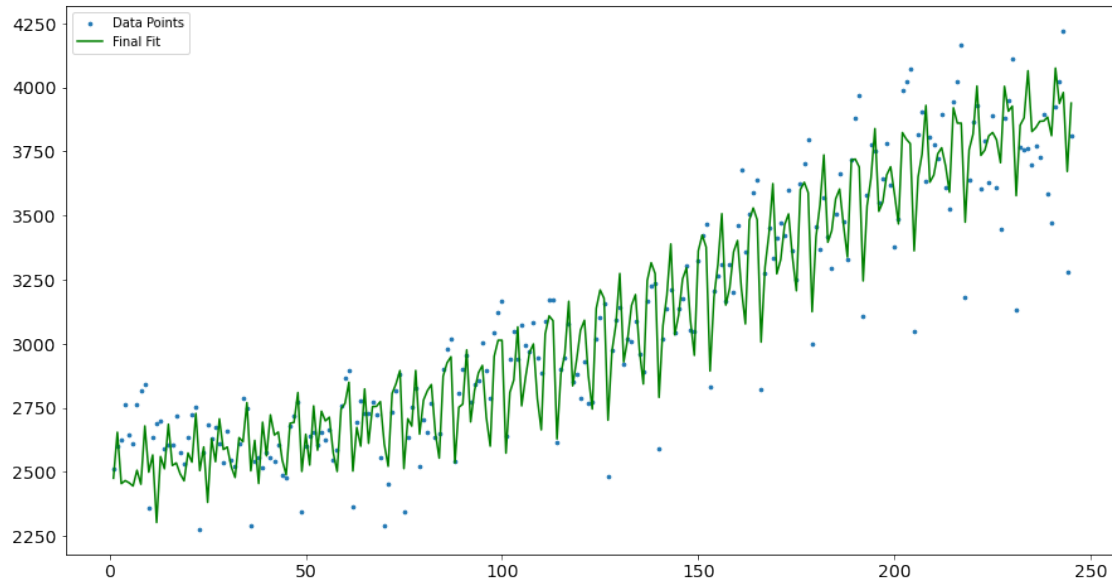
- Plot the combined fit together with a scatterplot of the data
- Compute and print the final **MSE** and R^2 coefficient. Comment on the improvement over the linear fit.

The combined fit, which corresponds to the full model

$$f(x) = \beta_0 + \beta_1 x + \beta_{s1} \sin(\omega_1 x) + \beta_{c1} \cos(\omega_1 x) + \dots + \beta_{sk} \sin(\omega_k x) + \beta_{ck} \cos(\omega_k x),$$

can be obtained by adding the fits in parts 2b) and 2e).

```
[14]: y_pred = fit + fit_sc
plt.scatter(xvals, yvals, s=7, label='Data Points')
plt.plot(xvals, y_pred, 'g', label='Final Fit')
plt.legend()
plt.show()
```



```
[15]: residual_total= yvals - y_pred
mean_total = residual_total.mean()
print(f'The mean of the new residuals is {round(mean_total,3)}')
MSE_final = residual_total.dot(residual_total) / N
R2_final  = 1- MSE_final/ yvals.var()
print('Mean squared error for the combined fit = ', np.round(MSE_final,3))
print('R^2 coefficient of combined fit = ', np.round(R2_final, 3))
```

The mean of the new residuals is 0.0

Mean squared error for the combined fit = 20297.501

R² coefficient of combined fit = 0.908

R² got improved from 0.8 to 0.9, while the MSE halved, so fitting sinusoidal components proved to model really well the seasonality of the data.